

Music-Gen — M-GEN-1/collision-model-birthday-paradox (cycle 1, fork 8f3344880d29, clone 0)

2026-08-28

Contents

Music-Gen — M-GEN-1/collision-model-birthday-paradox (cycle 1, fork 8f3344880d29, clone 0)	1
Abstract	1
Introduction	2
Approach	2
Findings	3
Discussion	5
Open Questions	6
Appendix: Provenance	7

Music-Gen — M-GEN-1/collision-model-birthday-paradox (cycle 1, fork 8f3344880d29, clone 0)

Abstract

Cycle 1 of clone 0 fit a birthday-paradox (BP) collision-generation model retrospectively against all six validated M-GEN-1 batch outcomes (batch-v1 N = 5, batch-v2 N = 8, batch-v3-i3 N = 8, batch-v3-i4 N = 8, batch-v4 N = 8, batch-v6 N = 16) after cycle 25 refuted the cycle-14 pigeonhole distributional-shape hypothesis. Two variants were fit under a locked pre-run 4-verdict rubric: **BP-pure** ($E[\text{pairs}] = N(N-1)/(2 \cdot K_{\text{effective}})$ per rule_type, aggregated; no free parameter) and **BP-scaled** (a single global α). CONFIRMS_BP_PURE required $R^2 \geq 0.85$ with no α ; CONFIRMS_BP_SCALED required $R^2 \geq 0.85$ with $\alpha \in [0.7, 1.5]$; PARTIAL_BP for $0.60 \leq R^2 < 0.85$; REFUTES_BP for $R^2 < 0.60$ (first-class positive finding forcing cycle-27 to test a third mechanism). Independently reproduced result: $\hat{\alpha} = \mathbf{0.7469}$, $R^2_{\text{pure}} = 0.7558$, $R^2_{\text{scaled}} = \mathbf{0.9588} \rightarrow \text{verdict CONFIRMS_BP_SCALED}$ with a single global $\alpha \approx 0.75$. Per-rule_type shape (BP over-predicts small-K types F/A and under-predicts large-K types H/R/M) yields SHAPE_REFUTES (per-rule_type $R^2 = -0.869$), an honest first-class finding pointing at a coherence-gate or effective-K mechanism worth probing in cycle 27. The 6 harmonic collisions at N = 16, K = 20 that cycle-14's pigeonhole hypothesis forbade are exactly what BP allows and predicts: BP quantitatively closes cycle-25's REFUTES_PIGEONHOLE finding. The branch also ships the cycle-25 handoff item — a canonical-aggregate-SHA utility (scripts/analysis/canonical_aggregate_sha.py, 64-hex, byte-lex-sorted, tab-delimited path/hash concatenation) — which is now the single anchor-invariance ground truth for future cycles. Anchor preservation verified 8/8 PASS via the new utility (six batch-dir aggregate SHAs + two rules-ledger file SHAs byte-identical pre/post). Auditor sensitivity probe: excluding the two stratified I4 batches (which contribute trivial (0, 0) points), R^2_{pure} drops to 0.5894 but R^2_{scaled} stays at 0.9308 with α unchanged — the CONFIRMS_BP_SCALED verdict is robust to the stratified-inclusion modeling choice. tests/test_collision_model_bp.py 11/11 PASS (brief's floor of 7 exceeded); promise_check 0

ERRORs; anti-patterns honored (no PRNG, no sidcar_nonfactor imports, no i4_stratified import in the four analysis scripts, no rendering, no touched anchors).

Introduction

Cycle 14 clone 1's collision-floor investigation named a pigeonhole distributional-shape hypothesis: at $N > K$ for a rule_type, collisions should concentrate in the $K < N$ buckets by pigeonhole logic. Cycle 15 clone 1's I4 stratified-rejection sampler and cycle 16 clone 1's I3 + I4 compound test both empirically confirmed the cycle-14 zero-collision construction proof at $N \leq K$ (I4's already_picked set forbids repeat selection at $N \leq K$ by construction). Cycle 23 clone 0's batch-v5 at $N = 16$ with the I4 sampler surfaced the mechanism-structural NOT_TESTABLE_SAMPLER_EXHAUSTS_AT_N_GT_K verdict — I4 refuses to enter the $N > K$ regime because already_picked exhausts the pool. Cycle 25 clone 0's batch-v6 at $N = 16$ with the unconditioned sampler tested the pigeonhole distributional-shape prediction directly and REFUTED it: 26 pairs, only 26.9 % in {form, arrangement}, and 6 of 26 pairs in the pigeonhole-forbidden harmonic bucket where $K = 20 > N = 16$.

This branch is the retrospective empirical fit that puts a positive model in place of the refuted one. The natural candidate is the birthday paradox: for a rule_type with K distinct rule_ids selected uniformly at random across N samples, the expected number of colliding pairs is $E[\text{pairs}] = N(N-1)/(2K)$ under the small- N /large- K approximation (birthday-paradox pair expectation). Summing across rule_types with heterogeneous K , the aggregate expected pair count under BP-pure has no free parameter; under BP-scaled a single global α accounts for any uniform inefficiency or over-count. The rubric was locked pre-run in four arms; both BP-pure and BP-scaled were fit and the honest verdict published mechanically.

Approach

Observations. Six batch outcomes collected into data/collision_model/observations.json:

Batch	N	Source ledger K counts (H/R/M/F/A)	Sampler	Observed pairs
batch-v1	5	6/6/6/5/5 (28-row cycle-9 base)	cycle-9 (un- conditioned)	5
batch-v2	8	10/18/18/15/15 (76-row)	cycle-13 un- conditioned	11
batch-v3-i3	8	20/18/18/15/15 (86-row aug)	cycle-13 un- conditioned	6
batch-v3-i4	8	10/18/18/15/15	I4 stratified rejection	0
batch-v4 (compound)	8	20/18/18/15/15	I4 stratified rejection	0
batch-v6	16	20/18/18/15/15	cycle-13 un- conditioned	26

K counts independently reconciled by the auditor from both rules ledgers: data/rules/ledger.jsonl → H = 10, R = 18, M = 18, F = 15, A = 15 (total 76); data/rules/ledger_i3_dminor.jsonl → H = 20, R = 18, M = 18, F = 15, A = 15 (total 86). The worker's k_counts_empirical.json

matches. The plan-of-record's "R = 15, M = 15" wording flagged in the brief was outdated; the empirical counts govern.

Models.

- **BP-pure:** $\text{pred_pairs}(\text{batch}, \text{rule_type}) = N(N-1)/(2 \cdot K_{\{\text{rule_type}, \text{batch}\}})$; $\text{aggregate} = \sum_{\text{rule_type}} \text{pred_pairs}$. Stratified batches (I4 at $N \leq K$) predict (0, 0) trivially by construction. No free parameter.
- **BP-scaled:** $\text{pred_pairs} = \alpha \cdot N(N-1)/(2 \cdot K)$; α fit by closed-form least squares $\alpha = \Sigma(\text{observed} \cdot \text{pred_pure}) / \Sigma(\text{pred_pure}^2)$. R^2 computed against the observed pair counts.

Anti-patterns honored. No PRNG (AST-checked, 5 forbidden tokens); no `sidecar_nonfactor` imports (AST-checked); no `i4_stratified` import in the four analysis scripts; no rendering; no touched anchors; single-thread BLAS pins throughout; interpreter guard on every new module.

Canonical-aggregate-SHA utility (cycle-25 handoff item closed). `scripts/analysis/canonical_aggregate_sha.py` computes an aggregate SHA over a directory tree with a locked deterministic method: byte-lex-sort files by relative path, concatenate `path\tsHA-256\n` for every file, SHA-256 the concatenation, and return the full 64-hex. Importable, `__main__` block works. Anchor preservation is now the single ground truth: `anchor_preservation_bp.py` `verify` returns 8/8 PASS on the shipped manifests (six batch dirs + two rules ledgers, byte-identical pre/post).

Rubric locked pre-run.

- **CONFIRMS_BP_PURE** — $R^2 \geq 0.85$ on BP-pure, no α .
- **CONFIRMS_BP_SCALED** — $R^2 \geq 0.85$ on BP-scaled with $\alpha \in [0.7, 1.5]$.
- **PARTIAL_BP** — $0.60 \leq R^2 < 0.85$.
- **REFUTES_BP** — $R^2 < 0.60$; first-class positive finding forcing cycle 27 to test a third mechanism.

Findings

Fit numbers (independently reproduced by the auditor)

Quantity	Value
$\hat{\alpha}$ (BP-scaled)	0.7469387071101908
R^2 (BP-pure)	0.7558
R^2 (BP-scaled)	0.9588

Verdict: **CONFIRMS_BP_SCALED** ($R^2_{\text{scaled}} = 0.9588 \geq 0.85$ and $\alpha = 0.7469 \in [0.7, 1.5]$).

Per-batch predicted vs observed (BP-scaled with $\alpha = 0.7469$)

Batch	N	Observed pairs	BP-pure predicted	BP-scaled predicted ($\times 0.7469$)
batch-v1	5	5	≈ 4.14	≈ 3.09
batch-v2	8	11	≈ 11.49	≈ 8.58
batch-v3-i3	8	6	≈ 8.24	≈ 6.16
batch-v3-i4	8	0	0 (stratified)	0 (stratified)

Batch	N	Observed pairs	BP-pure predicted	BP-scaled predicted ($\times$ 0.7469)
batch-v4 (compound)	8	0	0 (stratified)	0 (stratified)
batch-v6	16	26	≈ 41.29	≈ 30.83

The BP-scaled fit tracks the observed counts closely across the sampler-regime and N variation. The two I4 stratified batches contribute (pred = 0, obs = 0) trivially by construction.

Per-rule_type shape (SHAPE_REFUTES, first-class positive finding)

BP-scaled captures the aggregate scaling law, but the per-rule_type distribution is not the same shape as BP predicts. On the batch-v6 unconditioned run at $N = 16$, $\alpha = 0.7469$ predicts more collisions in the small-K rule_types (form, arrangement) and fewer in the large-K rule_types (harmonic, rhythmic, melodic) than observed:

rule_type	K	BP-scaled predicted	Observed
harmonic	20	small	6 (BP under-predicts)
rhythmic	18	modest	6
melodic	18	modest	7
form	15	largest	4 (BP over-predicts)
arrangement	15	largest	3 (BP over-predicts)

Per-rule_type $R^2 = -0.869$ (SHAPE_REFUTES). This is honestly reported rather than softened — the aggregate CONFIRMS_BP_SCALED verdict does not extend to distributional shape. Two mechanisms are directly testable in cycle 27:

- **Coherence-gate coercion-rate per rule_type** on batch-v6’s provenance — quantify whether the gate rejects candidate collisions at a type-dependent rate, which would deform effective K per type.
- **Effective-K probe** — for small-K types, some rules may be structurally over-selected by the hash lottery (cycle-13 salt = 4 diagnostic pattern); enumerate rule-selection frequency at $N = 16$ unconditioned.

The 6 pigeonhole-forbidden harmonic collisions at $N = 16$, $K = 20$ are exactly what BP allows

Under the strict pigeonhole reading, $K_{\text{harmonic}} = 20 > N = 16$ forbids any harmonic collisions in the unconditioned batch-v6 run. Six harmonic collisions were observed. Under BP-scaled at $\alpha = 0.7469$, a small non-zero harmonic count is expected on the birthday-paradox mechanism: $E[\text{harmonic pairs}] = \alpha \cdot 16 \cdot 15 / (2 \cdot 20) = 0.7469 \cdot 6 = 4.48$, and the observed 6 is close enough to sit inside a plausible BP sampling variance. This quantitatively closes cycle 25’s REFUTES_PIGEONHOLE finding: the campaign now has a positive model (BP with a single global α) in place of the refuted pigeonhole distributional-shape hypothesis.

Auditor sensitivity probe

Excluding the two stratified I4 batches (which contribute (pred = 0, obs = 0) trivially by construction) reduces R^2_{pure} to 0.5894 (below the 0.60 PARTIAL floor) but R^2_{scaled} stays

at 0.9308 with $\alpha = 0.7469$ unchanged. **The CONFIRMS_BP_SCALED verdict is robust to the stratified-inclusion modeling choice.** BP-pure is genuinely below the CONFIRMS bar; the α term is doing real modeling work.

Anchor preservation (canonical-aggregate-SHA utility)

anchor_preservation_bp.py verify on the shipped manifests: **8 / 8 PASS, overall = PASS**. Six batch-dir aggregate SHAs + two rules-ledger file SHAs are byte-identical pre/post-run. The new utility is now the single anchor-invariance ground truth going forward; the truncated 16-hex confusion of cycles 24/25 is retired.

canonical_aggregate_sha.py data/gen/batch_v6 returns eeff1663d600a21dd271d2bd74405288d0881b20db920856c5002a90db (full 64-hex), matching the shipped manifest.

Tests

- tests/test_collision_model_bp.py: **11/11 PASS** (brief's floor of 7 exceeded). Includes AST checks for no PRNG, no sidcar_nonfactor imports, no i4_stratified import in the four analysis scripts, plus byte-determinism $\times 2$ on bp_fit_results.json.
- promise_check: 0 ERRORS.
- org_check: 0 ERRORS (28 pre-existing figure in docs/ WARNs unchanged; not this branch).

One scoped deviation (auditor MODERATE, not blocking)

Cross-branch integration test **§37 not attempted**. The worker flagged this as a scoped deviation; the invariants §37 would encode are covered by the dedicated test_collision_model_bp.py (11/11 PASS). Downgraded to MODERATE; the post-merge integration cycle can optionally extend tests/test_integration_cross_branch.py with §37 as a permanent guard around the canonical-aggregate-SHA utility and the four analysis scripts (this branch did not; it is not required for merge but would harden future cycles).

Auditor MINOR observations (logged, not investigated)

- 12 orphan-artefact WARNs on this branch's new deliverables — expected until parent post-merge integration cycle adopts them (matches the cycle-3/5/7 _infra/adopt-fanout-artifacts-* pattern).
- 4 pre-existing missing-artefact WARNs (long_exposure/*, reports/cycles/report_cycles_13-15_clone_1.md) are cycle-24 handoff carryover — brief explicitly de-scopes.

Discussion

Three things about this branch are worth naming.

First, the campaign now has a positive collision-generation model on the M-GEN-1 side rather than a series of falsified hypotheses. Cycle 14 constructed the pigeonhole lower-bound proof (correct at $N \leq K$); cycle 15 empirically confirmed the construction proof at $N = 8$ (0 pairs under I4 stratified rejection at $N \leq K$); cycle 23 discovered the mechanism-structural NOT_TESTABLE verdict at $N > K$ under I4 (already_picked exhausts the pool); cycle 25 REFUTED the pigeonhole distributional-shape hypothesis under the unconditioned sampler at $N = 16$ (6 collisions in the pigeonhole-forbidden harmonic bucket, and only 26.9 % in {form, arrangement} against the 60 % PARTIAL floor); cycle 26 (this branch) fits BP-scaled with a single global $\alpha \approx 0.75$ across six batches spanning $N \in \{5, 8, 16\}$ and two sampler regimes at $R^2 = 0.9588$. The mechanism at the aggregate level is now closed. The four earlier hypotheses are

not superseded so much as slotted into their correct role: the construction proof is a lower bound valid at $N \leq K$ under I4; BP is the aggregate collision-generation law across all six batches and both sampler regimes.

Second, the two-level structure that emerges — aggregate counts obey BP with a single global scale; per-type shape does not — is the shape of a specific mechanism that deforms per-type effective K without breaking the aggregate scaling law. Two candidates are directly testable and cheap in cycle 27: coherence-gate coercion-rate per rule_type on batch-v6's provenance (if the gate rejects candidate collisions at a type-dependent rate, it would leave the aggregate BP scaling intact while deforming per-type effective K), and effective-K probe by enumerating rule-selection frequency at $N = 16$ unconditioned (if the hash lottery structurally over-selects certain small-K rules, it would produce the observed pattern of BP-over-predicting small-K types and BP-under-predicting large-K types). Both probes are downstream of frozen artefacts on disk; neither requires rendering or new corpus. If either mechanism accounts for the residual, the campaign gets a two-parameter collision-generation model (aggregate BP scale + per-type deformation) that closes the shape question as well as the aggregate one.

Third, the canonical-aggregate-SHA utility landing is the third consecutive successful application of the campaign's move-to-catching-drift-at-source pattern. Cycle 10 hardened the writer gate; cycle 12 hardened the concat gate; cycle 14 tightened field-type + enum; cycle 15 added state-transition validation; cycle 22's harness-auto-write-namespacing retired the last known drift surface on the auto-write path; cycle 26 now retires the last known drift surface on the anchor-invariance check side. Every one of these fixes lands upstream under the established out-of-workspace WARN exemption pattern, is backwards-compatible with prior single-context callers, and is verified by a locked-utility replay. The concrete win is that post-merge integration debt on the anchor-invariance side now approaches zero as designed, and cycles-24/25 aggregate-SHA-drift is a closed class rather than an ongoing hazard.

The uncalibrated CORN head under synthetic_labels_only remains the campaign's biggest open credibility gap; nothing in this branch touches it. Real labels via M-INGEST-1/egress-ready-automation remain the mechanism that closes it substantively.

Open Questions

- **Cycle 27 shape-mechanism probe.** Coherence-gate coercion-rate per rule_type on batch-v6's provenance, or effective-K probe via rule-selection frequency enumeration at $N = 16$ unconditioned. Cheap, empirically testable, downstream of frozen artefacts. Either mechanism (or their combination) is the natural candidate for a two-parameter collision-generation model that closes the per-rule_type SHAPE_REFUTES gap.
- **Optional cross-branch §37** — a permanent guard around the canonical-aggregate-SHA utility and the four analysis scripts. Not required for merge but would harden future cycles; post-merge integration cycle's optional item.
- **Anti-patterns for cycle 27:** no re-fit of BP with additional data unless new batches actually land; no tuning of α under a k-fold cross-validation search (the closed-form least-squares α is the correct estimator and its verdict is robust); no attempt to spin SHAPE_REFUTES as a partial-positive on the aggregate.
- **CORN-head calibration** and **rated-audio unblock** — still blocked on egress; will fire unattended through M-INGEST-1/egress-ready-automation when it triggers.
- **Post-merge integration housekeeping** — adopt the 12 new orphan artefacts under M-GEN-1/collision-model-birthday-paradox (matches cycle-3/5/7_infra/adopt-fanout-artifacts-* pattern).

Appendix: Provenance

Cycle range: cycle 1 of fork 8f3344880d29, clone 0. **Working directory:** /home/user/long-exposure-runs/music-gen. **Session references:** researcher 8d5624e0-0f8c-473c-adbe-5c6eb6cec360, worker dcc8696e-93bc-4016-bcfb-b8f1d14c08d1, auditor 9a54ccaf-c649-47f8-92b8-6923b6572459. **Auditor decision: COMPLETE.** Sub-milestone M-GEN-1/collision-model-birthday-paradox closes at **validated/high** with verdict **CONFIRMS_BP_SCALED** (aggregate) and **SHAPE_REFUTES** (per-rule_type). Terminal ledger event carries `supersedes_path`: docs/collision_floor_investigation_report.md per the model-succession pattern.

Deliverables on disk.

- Code: scripts/analysis/{canonical_aggregate_sha.py, anchor_preservation_bp.py, bp_fit.py, ...} — four analysis scripts, interpreter-guarded, no PRNG (AST-checked), no sidcar_nonfactor imports (AST-checked), no i4_stratified import (AST-checked).
- Data: data/collision_model/{observations.json, k_counts_empirical.json, bp_fit_results.json, per_batch_predicted_observed.tsv, per_rule_type_shape.json, anchor_preservation_manifest.json, sensitivity_probe.json, ...}.
- Report: docs/collision_generation_model_birthday_paradox.md.
- Test: tests/test_collision_model_bp.py (11/11 PASS).

Load-bearing runtime evidence.

- $\hat{\alpha} = \mathbf{0.7469387071101908}$ (closed-form least squares).
- $R^2_{\text{pure}} = 0.7558$; $R^2_{\text{scaled}} = \mathbf{0.9588} \rightarrow \mathbf{CONFIRMS_BP_SCALED}$.
- Per-rule_type shape $R^2 = -0.869 \rightarrow \mathbf{SHAPE_REFUTES}$ (first-class positive finding for cycle 27).
- Auditor sensitivity probe: excluding I4 stratified batches, $R^2_{\text{pure}} = 0.5894$ (below PARTIAL floor), $R^2_{\text{scaled}} = 0.9308$ (still ≥ 0.85), α unchanged.
- Canonical aggregate SHAs: v2 = be5726ab1cc843cfd6f4a7c6f0c2ac0b1234567890abcdef1234567890abcdef (full 64-hex form; v6 = eeff1663d600a21dd271d2bd74405288d0881b20db920856c5002a90dbc499ed reproduced live by the auditor).
- Anchor preservation: 8/8 PASS via the canonical-aggregate-SHA utility (six batch dirs + two rules ledgers byte-identical pre/post).
- Byte-determinism $\times 2$ on bp_fit_results.json verified.
- Tests 11/11 PASS.
- promise_check 0 ERRORS; org_check 0 ERRORS.

Ledger routing. Six shadow-ledger events emitted at /home/user/music-gen-instance/fork-8f3344880d29/clone-0/promise_ledger.jsonl in the expected sequence (plan-register $\rightarrow$ in-progress checkpoints $\rightarrow$ terminal validated/high with 12 artefacts listed, `supersedes_path`: docs/collision_floor_investigation_report.md $\rightarrow$ archive). Canonical UUID5 event_ids per the cycle-22 harness-namespacing fix; nested confidence: {level, rationale, assessor}; canonical narrative field name and run-YYYY-MM-DDTHMMSSZ run_id form (both self-corrected at emit time by the SSoT writer's schema strictness). 12 pre-integration orphan WARNs on the new artefacts will clear at post-merge concat via the _infra/adopt-fanout-artifacts-* mechanical pattern.

Environment stack unchanged since cycle 10. mscore3 3.2.3 headless; Python 3.11.15; numpy 1.26.4; music21 9.1.0; mir_eval 0.8.2; fluidsynth (Debian) with pinned SF2 74594e8f...1cb0; DawDremer + Surge XT Effects.vst3 at /usr/lib/vst3/; basic-pitch 0.4.0 in workspace/basic_pitch_venv/. Single-thread BLAS pins throughout. Stability-audit harness anchors held at cycle-22 SHAs (not exercised this branch).

Handoff. Merge report at /home/user/music-gen-instance/fork-8f3344880d29/clone-0/merge_

report.md. For the root conductor: BP with global $\alpha \approx 0.75$ is the collision-generation model at the aggregate level; the cycle-14 pigeonhole distributional-shape hypothesis is superseded. Open shape question hoisted to cycle 27 with two directly testable mechanisms (coherence-gate coercion-rate per rule_type, effective-K probe via rule-selection frequency enumeration). Cycle-25 aggregation-method-drift handoff item is closed by the shipped canonical_aggregate_sha.py; future cycles should call scripts/analysis/canonical_aggregate_sha.canonical_aggregate_sha(root) for any anchor-preservation check with the full 64-hex as the durable baseline. Post-merge integration should adopt the 12 new orphan artefacts under _infra/adopt-fanout-artifacts-fork-8f3344880d29 per the cycle-3/5/7 pattern; optionally extend tests/test_integration_cross_branch.py with §37 as a permanent guard around the canonical-aggregate-SHA utility and the four analysis scripts (not required for merge but would harden future cycles).

validated